Machine Learning for Customer Churn Prediction

Predicting why customers left the company



OSEMN Process | Models Used

Obtain K Nearest Neighbors Classifier

Scrub Gradient Boosting Classifier

Explore Random Forest Classifier

Model Support Vector Classifier

Interpret Adaboost Classifier

Telecom Customer Churn Data

The data used was from a telecom company, with information like `phone number`, `day minutes used`, and length an account has been established.

The data was downloaded from kaggle under the title `Churn in Telecom dataset`



Scrubbing the Data:

There were no null values in the data, only some yes/nos that needed to be encoded into 1/0 respectively.

Encoding:

We do this because a 0 or a 1 is easier for a computer to understand vs a yes or a no.

Questions to Answer

What geographical locations should the company focus on?

What type of model will perform best on the given dataset?

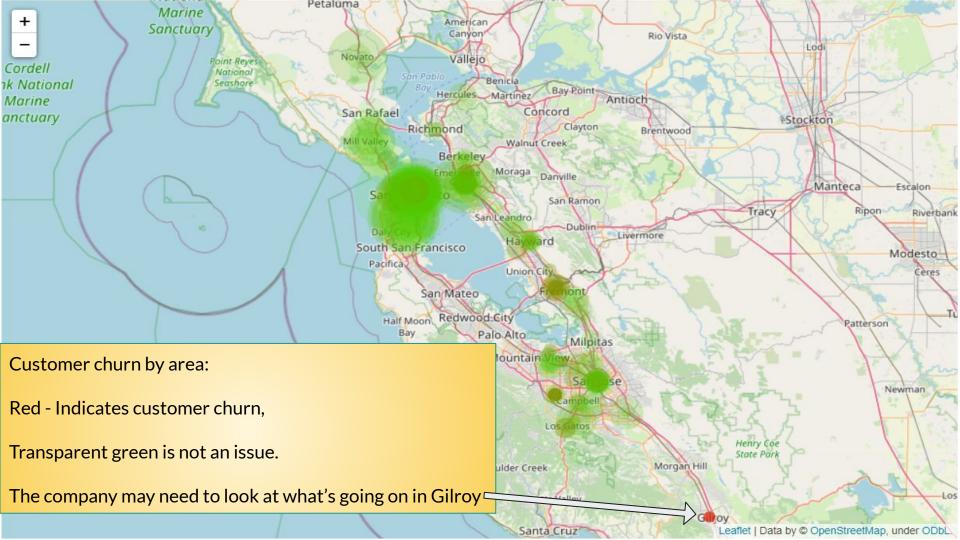
Why does a customer decide to leave their current provider?

How can the company improve their customer retention?

Explore:

We used phone number for location scraping the for relative longitude and latitude of a customer for the map ahead.

Note: This could be skewed by someone moving, but keeping their phone number, as in the data we have multiple states, but all of the area codes are from California.



Modeling:

After trying the models outlined in the first slide we found that our Gradient Boosting Classifier performed the best.

Gradient Boosting Classifier:

We used a sklearn's GridSearchCV to find that the most accurate parameters were a learning rate of one and a max depth of ten. Giving us a overall testing score of ~96%, and ~68% of the time being able to say whether a customer would leave or not.

Don't be confused by the squares on the right:

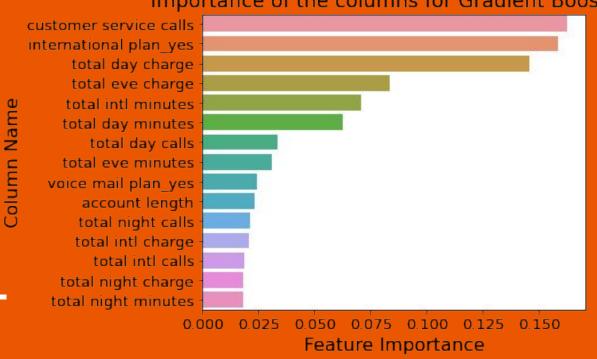
Bottom right is how many customers leaving our model successfully predicted.

Top left is how many customers staying our model successfully predicted.



Feature Importance:

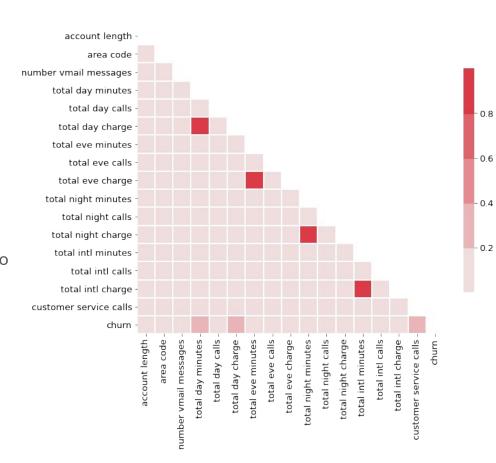




Feature Correlation:

This is a pay per minute service as total minutes, and total charge are strongly correlated, which explains why customer churn goes up as minutes go up their bill is higher.

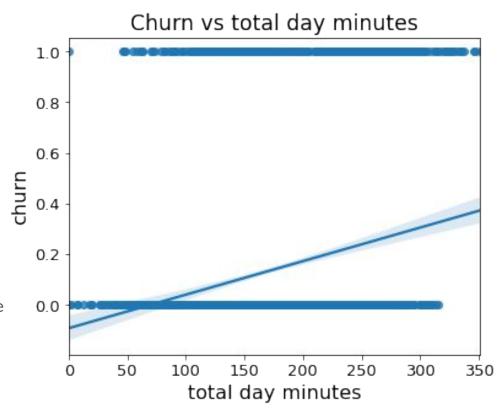
Total day minutes, charge, and customer service calls all slightly correlate with churn. As shown in the feature importances.



Total Day Minutes:

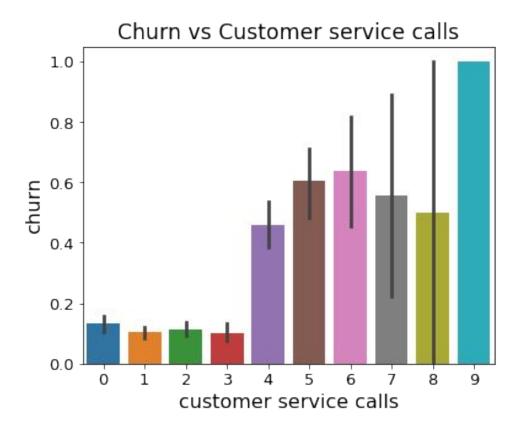
The graph is showing that the more minutes a customer uses the more likely they are to leave the provider.

As total day minutes correlates to total day charge the company should think about lowering their rates.



Customer Service Calls:

It seems that the more customer service calls a customer is making, the more likely they are to leave the company or "churn".



Conclusion:

- The company should focus on their customer service calls, and how much they are charging for their minutes.
- Out of the models we tried Gradient Boosting performed the best on the given data.
- The company can improve their customer retention by changing their customer service center, and lowering the price of day minutes.

Future Work:

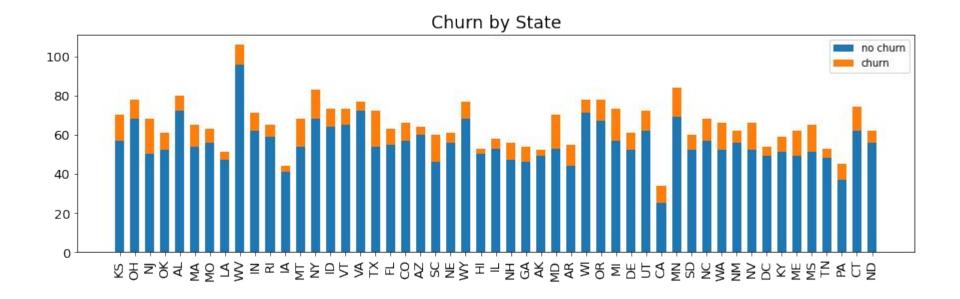
- Spending more processing power tuning the models, as all the models ran in under 30 seconds.
- Try different models like XGBoost, Extra Trees, and Logistic Classification.
- Running a Linear Regression on some of the important features to extract coefficients.

Thank you.



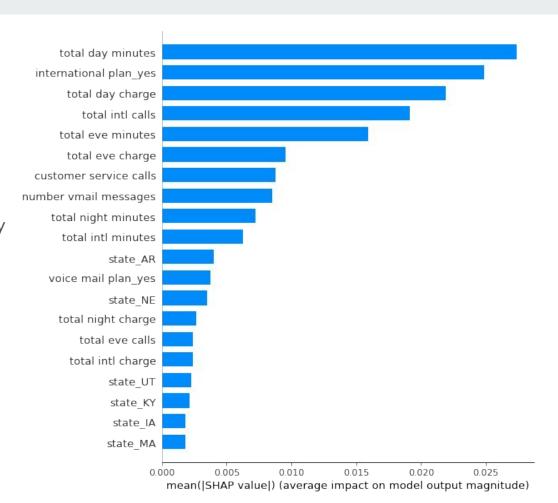


Churn Distribution:



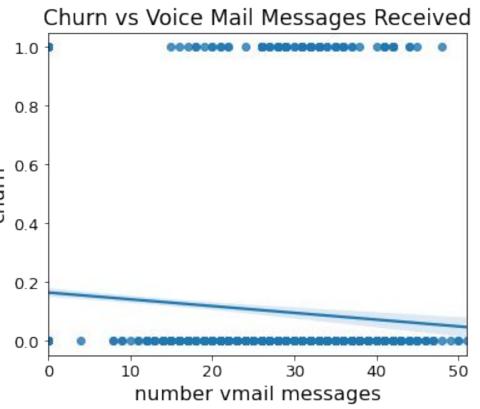
Shap feature importance

Just like the other importance chart day minutes and international plan are the most important. A key difference from the other importance is customer service calls is not at the top.



Voice Mails

Maybe the company should leave more voice mails if they are not currently, as the more voice mails someone receives, the more likely they are to stay with the company.



Voicemail Distribution:

On average, people receive more voicemails than I do.

